

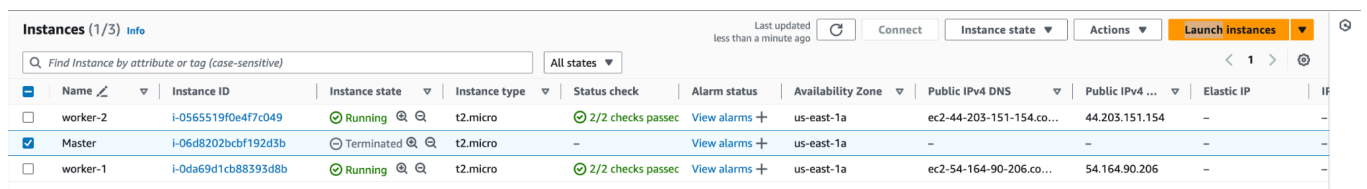
Advanced DevOps Lab

Experiment:3

Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Steps:

1. Create 3 EC2 Ubuntu Instances on AWS.

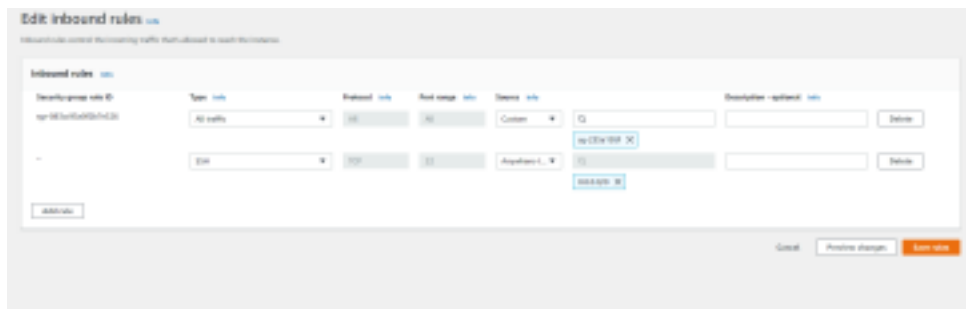


The screenshot shows the AWS Management Console 'Instances' page. It lists three EC2 instances in the us-east-1 region. The 'Master' instance is selected and highlighted in blue. It is a t2.micro instance with a 'Running' status, 2/2 checks passed, and an alarm status of 'View alarms'. The other two instances, 'worker-2' and 'worker-1', are also t2.micro instances with a 'Running' status and 2/2 checks passed. The 'Launch Instances' button is visible in the top right corner.

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...	Elastic IP
worker-2	i-0565519f0e4f7c049	Running	t2.micro	2/2 checks passed	View alarms	us-east-1a	ec2-44-203-151-154.co...	44.203.151.154	-
Master	i-06d8202bcbf192d3b	Running	t2.micro	2/2 checks passed	View alarms	us-east-1a	-	-	-
worker-1	i-0da69d1cb88393d8b	Running	t2.micro	2/2 checks passed	View alarms	us-east-1a	ec2-54-164-90-206.co...	54.164.90.206	-

(Name 1 as Master, the other 2 as worker-1 and worker-2)

-2. Edit the Security Group Inbound Rules to allow SSH



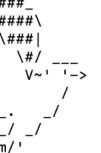
The screenshot shows the 'Edit inbound rules' page for a security group. It displays a table of inbound rules. The first rule is 'All traffic' with a type of 'All traffic', protocol of 'All', and port range of 'All'. The second rule is 'SSH' with a type of 'SSH', protocol of 'TCP', and port range of '22'. The 'SSH' rule is highlighted in blue. The 'Add rule' button is visible at the bottom right.

Security group rule ID	Type	Protocol	Port range	Source	Description
sg-1a1b1c1d	All traffic	All	All	Custom	
sg-2a2b2c2d	SSH	TCP	22	Anywhere	

3. SSH into all 3 machines

ssh -i <keyname>.pem ubuntu@<public_ip_address>

```
(base) krushikeshunilshelar@Krushikesh-S MacBook-Air downloads % ssh -i "newkey.pem" ec2-user@ec2-44-203-151-154.compute-1.amazonaws.com
The authenticity of host 'ec2-44-203-151-154.compute-1.amazonaws.com ([44.203.151.154])' can't be established.
ED25519 key fingerprint is SHA256:4cidwEvWyyoqWE0gGMSqDMjX2SlxkZVTUTIbdzMDdlic.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-44-203-151-154.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
```



```
#_
~\#####      Amazon Linux 2023
~~~\#####\
~~~~\####|
~~~~\#/____https://aws.amazon.com/linux/amazon-linux-2023
~~~~V~'|>
~~~~
~~~~_. _/_/_/
~~~~_/ /'_/_/_/
~~~~/_m/'
```

```
[ec2-user@ip-172-31-95-91 ~]$ ssh -i "newkey.pem" ec2-user@ec2-54-164-90-206.compute-1.amazonaws.com
Warning: Identity file newkey.pem not accessible: No such file or directory.
The authenticity of host 'ec2-54-164-90-206.compute-1.amazonaws.com ([172.31.88.50])' can't be established.
ED25519 key fingerprint is SHA256:RRoSz1NvNg9JLCAJDhKUUn6FRcRu1vtNbkiJV05M/I.
This key is not known by any other names
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-54-164-90-206.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
ec2-user@ec2-54-164-90-206.compute-1.amazonaws.com: Permission denied (publickey,gssapi-keyex,gssapi-with-mic).
[ec2-user@ip-172-31-95-91 ~]$
```

4. From now on, until mentioned, perform these steps on all 3 machines.
- ```
sudo yum install docker -y
```

```
[ec2-user@ip-172-31-92-18 ~]$ sudo yum install docker -y
Last metadata expiration check: 0:09:56 ago on Wed Sep 11 15:19:39 2024.
Dependencies resolved.
```

| Package                | Architecture |
|------------------------|--------------|
| docker                 | x86_64       |
| containerd             | x86_64       |
| iptables-libs          | x86_64       |
| iptables-nft           | x86_64       |
| libcgroup              | x86_64       |
| libnetfilter_conntrack | x86_64       |
| libnfnl                | x86_64       |
| libnftnl               | x86_64       |
| pigz                   | x86_64       |
| runc                   | x86_64       |

```
Transaction Summary
```

Then, configure cgroup in a daemon.json file by using following commands

- `cd /etc/docker`
- `cat <<EOF | sudo tee /etc/docker/daemon.json`  
{  
"exec-opts": ["native.cgroupdriver=systemd"],  
"log-driver": "json-file",

```
"log-opts": {
 "max-size": "100m"
},
"storage-driver": "overlay2"
}
EOF
```

- `sudo systemctl enable docker`
- `sudo systemctl daemon-reload`
- `sudo systemctl restart docker`
- `docker -v`

## Install Kubernetes on all 3 machines

SELinux needs to be disabled before configuring kubelet

- `sudo setenforce 0`
- `sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config`

```
[ec2-user@ip-172-31-81-63 docker]$ sudo setenforce 0
[ec2-user@ip-172-31-81-63 docker]$ sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

Add kubernetes repository (paste in terminal)

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd
.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
```

Type following commands:

- `sudo yum update`

- `sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes`

```
[ec2-user@ip-172-31-81-63 docker]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Last metadata expiration check: 0:01:34 ago on Wed Sep 11 15:39:05 2024.
Dependencies resolved.
=====
Package Architecture Version
=====
Installing:
kubeadm x86_64 1.30.4-150500.1.1
kubectl x86_64 1.30.4-150500.1.1
kubelet x86_64 1.30.4-150500.1.1
Installing dependencies:
conntrack-tools x86_64 1.4.6-2.amzn2023.0.2
cri-tools x86_64 1.30.1-150500.1.1
kubernetes-cni x86_64 1.4.0-150500.1.1
libnetfilter_cthelper x86_64 1.0.0-21.amzn2023.0.2
libnetfilter_cttimeout x86_64 1.0.0-19.amzn2023.0.2
libnetfilter_queue x86_64 1.0.5-2.amzn2023.0.2
socat x86_64 1.7.4.2-1.amzn2023.0.2
Transaction Summary
=====
Install 10 Packages
```

After installing Kubernetes, we need to configure internet options to allow bridging.

- `sudo swapoff -a`
- `echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf`
- `sudo sysctl -p`

## 1. Perform this ONLY on the Master machine

Initialize kubernetes by typing below command

- `sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all`

```
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

 mkdir -p $HOME/.kube
 sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
 sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

 export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
 https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.81.63:6443 --token zh5jbb.a6ty3eujzc51d15d \
 --discovery-token-ca-cert-hash sha256:0822f656bf52a17a2b6686c123f811306f41495ca650a0aed9bf6cd2d2f6f8c5
[ec2-user@ip-172-31-81-63 docker]$ mkdir -p $HOME/.kube
 sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
 sudo chown $(id -u):$(id -g) $HOME/.kube/config
[ec2-user@ip-172-31-81-63 docker]$
```

Copy the mkdir and chown commands from the top and execute them

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

**Copy this join link and save it in clipboard (copy from your output as it different for each instance)**

My personal join key:

```
kubeadm join 172.31.92.157:6443 --token x4sw6q.sbckmhm5gkoubquv \
 --discovery-token-ca-cert-hash
sha256:24c005691fcab2260667ee43384d46afd4b2b27401e82c01550798a0d8f98950
```

Then, add a common networking plugin called flannel file as mentioned in the code.

**This step gives error:**

```
kubectl apply -f
```

```
https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```